

## ✓ 1.Import Required Libraries

```
# Importing necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

## ✓ 2.Load the Dataset

```
from google.colab import files
uploaded= files.upload()
```




No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving Telco-Customer-Churn.csv to Telco-Customer-Churn.csv

```
df = pd.read_csv("Telco-Customer-Churn.csv")
print(df.head())
```



	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	\
0	7590-VHVEG	Female	0	Yes	No	1	No	
1	5575-GNVDE	Male	0	No	No	34	Yes	
2	3668-QPYBK	Male	0	No	No	2	Yes	
3	7795-CFOCW	Male	0	No	No	45	No	
4	9237-HQITU	Female	0	No	No	2	Yes	

	MultipleLines	InternetService	OnlineSecurity	...	DeviceProtection	\
0	No phone service	DSL	No	...	No	
1	No	DSL	Yes	...	Yes	
2	No	DSL	Yes	...	No	
3	No phone service	DSL	Yes	...	Yes	
4	No	Fiber optic	No	...	No	

	TechSupport	StreamingTV	StreamingMovies	Contract	PaperlessBilling	\
0	No	No	No	Month-to-month	Yes	
1	No	No	No	One year	No	
2	No	No	No	Month-to-month	Yes	
3	Yes	No	No	One year	No	
4	No	No	No	Month-to-month	Yes	

	PaymentMethod	MonthlyCharges	TotalCharges	Churn
0	Electronic check	29.85	29.85	No
1	Mailed check	56.95	1889.5	No
2	Mailed check	53.85	108.15	Yes
3	Bank transfer (automatic)	42.30	1840.75	No
4	Electronic check	70.70	151.65	Yes

[5 rows x 21 columns]

## ✓ Data Preprocessing

```
# Check for missing values
```

```
print(df.isnull().sum())
```

```
# Convert TotalCharges to numeric (it has missing/blank values)
```

```
df["TotalCharges"] = pd.to_numeric(df["TotalCharges"], errors='coerce')
```

```
# Fill missing values with median
```

```
df["TotalCharges"].fillna(df["TotalCharges"].median(), inplace=True)
```



```
customerID      0
gender          0
SeniorCitizen   0
Partner         0
Dependents      0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    0
Churn           0
```

```
dtype: int64
```

<ipython-input-8-29992b9cd5e8>:8: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series consisting of rows and columns but using inplace=True will not have any effect as a new object will be created. The behavior will change in pandas 3.0. This inplace method will never work because the

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col

```
df["TotalCharges"].fillna(df["TotalCharges"].median(), inplace=True)
```

```
# Check for duplicates
print("Duplicates:", df.duplicated().sum())

# Drop duplicates if any
df.drop_duplicates(inplace=True)
```

⇒ Duplicates: 0

```
import seaborn as sns
import matplotlib.pyplot as plt
```


```
# Boxplot before removing outliers
sns.boxplot(x=df["MonthlyCharges"])
plt.title("Before Removing Outliers - MonthlyCharges")
plt.show()
```

```
# Remove outliers using IQR
Q1 = df["MonthlyCharges"].quantile(0.25)
Q3 = df["MonthlyCharges"].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
df = df[(df["MonthlyCharges"] >= lower_bound) & (df["MonthlyCharges"] <= upper_bound)]
```


```
# Boxplot after removing outliers
sns.boxplot(x=df["MonthlyCharges"])
plt.title("After Removing Outliers - MonthlyCharges")
plt.show()
```



Before Removing Outliers - MonthlyCharges



After Removing Outliers - MonthlyCharges



```

from sklearn.preprocessing import LabelEncoder

df_encoded = df.copy()
categorical_cols = df_encoded.select_dtypes(include='object').columns.tolist() # Convert to

# Drop customerID as it's not useful
if 'customerID' in categorical_cols: # Check if 'customerID' is in the list before removing
    categorical_cols.remove('customerID')
    df_encoded.drop("customerID", axis=1, inplace=True)

# Label encode binary categorical features
le = LabelEncoder()
for col in categorical_cols:
    # Ensure the column still exists in df_encoded after potential drops
    if col in df_encoded.columns:
        if df_encoded[col].nunique() == 2:
            df_encoded[col] = le.fit_transform(df_encoded[col])
        # The 'elif col != 'customerID':` is no longer needed since 'customerID' is removed
    else:
        df_encoded = pd.get_dummies(df_encoded, columns=[col], prefix=col, drop_first=True)

print(df_encoded.head())

```

```

➡ gender SeniorCitizen Partner Dependents tenure PhoneService \
0      0      0      1      0      1      0
1      1      0      0      0      34      1
2      1      0      0      0      2      1
3      1      0      0      0      45      0
4      0      0      0      0      2      1

PaperlessBilling MonthlyCharges TotalCharges Churn ... \
0      1      29.85      29.85      0 ...
1      0      56.95     1889.50      0 ...
2      1      53.85      108.15      1 ...
3      0      42.30     1840.75      0 ...
4      1      70.70     151.65      1 ...

TechSupport_Yes StreamingTV_No internet service StreamingTV_Yes \
0      False      False      False      False
1      False      False      False      False
2      False      False      False      False
3      True       False      False      False
4      False      False      False      False

StreamingMovies_No internet service StreamingMovies_Yes \
0      False      False
1      False      False
2      False      False
3      False      False
4      False      False

```

```

Contract_One year Contract_Two year \

```

0	False	False
1	True	False
2	False	False
3	True	False
4	False	False

	PaymentMethod_Credit card (automatic)	PaymentMethod_Electronic check \
0	False	True
1	False	False
2	False	False
3	False	False
4	False	True

	PaymentMethod_Mailed check
0	False
1	True
2	True
3	False
4	False

[5 rows x 31 columns]

```
from sklearn.preprocessing import StandardScaler
```

```
# Scale numerical columns
```

```
scaler = StandardScaler()
```

```
numeric_cols = ['tenure', 'MonthlyCharges', 'TotalCharges']
```

```
df_encoded[numeric_cols] = scaler.fit_transform(df_encoded[numeric_cols])
```

```
print(df_encoded[numeric_cols].describe())
```

```

tenure  MonthlyCharges  TotalCharges
count    7.043000e+03    7.043000e+03    7.043000e+03
mean    -2.421273e-17    -6.406285e-17    -1.488074e-17
std      1.000071e+00     1.000071e+00     1.000071e+00
min     -1.318165e+00    -1.545860e+00    -9.991203e-01
25%     -9.516817e-01    -9.725399e-01    -8.298459e-01
50%     -1.372744e-01     1.857327e-01    -3.904632e-01
75%      9.214551e-01     8.338335e-01     6.642871e-01
max      1.613701e+00     1.794352e+00     2.826743e+00

```


```
sns.histplot(df["MonthlyCharges"], kde=True)
```

```
plt.title("Before Scaling - MonthlyCharges")
```

```
plt.show()
```



Before Scaling - MonthlyCharges



```
# Assuming you want to scale numerical features and store the result in df_encoded
# This is a placeholder and might need adjustment based on your full data processing steps
```

```
from sklearn.preprocessing import StandardScaler
```


```
# Select the numerical column(s) to scale
numerical_cols = ['MonthlyCharges']
```

```
# Create a copy of the original DataFrame to avoid modifying it directly
df_encoded = df.copy()
```

```
# Initialize the StandardScaler
scaler = StandardScaler()
```

```
# Fit and transform the numerical column(s)
df_encoded[numerical_cols] = scaler.fit_transform(df_encoded[numerical_cols])
```

```
# Now you can plot the scaled 'MonthlyCharges'
sns.histplot(df_encoded["MonthlyCharges"], kde=True)
plt.title("After Scaling - MonthlyCharges")
plt.show()
```



## ✓ EDA

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Load dataset
df = pd.read_csv("Telco-Customer-Churn.csv")

# Convert TotalCharges to numeric and handle missing values
df["TotalCharges"] = pd.to_numeric(df["TotalCharges"], errors='coerce')
df["TotalCharges"].fillna(df["TotalCharges"].median(), inplace=True)

# Drop customerID
df.drop("customerID", axis=1, inplace=True)

# Encode target variable
df["Churn"] = df["Churn"].map({"Yes": 1, "No": 0})

# Set seaborn style
sns.set(style="whitegrid")

# Histograms
```



```
df[["tenure", "MonthlyCharges", "TotalCharges"]].hist(bins=30, figsize=(10, 6), color='skybl  
plt.suptitle("Histograms of Numerical Features")  
plt.show()  
  
# Boxplot: Monthly Charges vs Churn  
plt.figure(figsize=(10, 4))  
sns.boxplot(x='Churn', y='MonthlyCharges', data=df)  
plt.title('Monthly Charges vs Churn')  
plt.show()
```



## ✓ Model Building

```
# Import required libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import seaborn as sns
import matplotlib.pyplot as plt

# Load dataset
df = pd.read_csv("Telco-Customer-Churn.csv")

# Data Preprocessing
df.drop('customerID', axis=1, inplace=True)
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], errors='coerce')
df['TotalCharges'].fillna(df['TotalCharges'].median(), inplace=True)


# Encode categorical variables
le = LabelEncoder()
df['Churn'] = le.fit_transform(df['Churn']) # Yes/No to 1/0
for column in df.select_dtypes(include='object').columns:
    if df[column].nunique() == 2:
        df[column] = le.fit_transform(df[column])
    else:
        df = pd.get_dummies(df, columns=[column])

# Feature Scaling
scaler = StandardScaler()
df[['tenure', 'MonthlyCharges', 'TotalCharges']] = scaler.fit_transform(df[['tenure', 'MonthlyCharges', 'TotalCharges']])


# Split data
X = df.drop('Churn', axis=1)
y = df['Churn']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# -----
# Logistic Regression
```

```
# -----  
log_model = LogisticRegression(max_iter=1000)  
log_model.fit(X_train, y_train)  
y_pred_log = log_model.predict(X_test)  
  
print("📊 Logistic Regression Results:")  
print("Accuracy:", accuracy_score(y_test, y_pred_log))  
print(classification_report(y_test, y_pred_log))  
sns.heatmap(confusion_matrix(y_test, y_pred_log), annot=True, fmt='d', cmap='Blues')  
plt.title("Logistic Regression - Confusion Matrix")  
plt.show()  
  
# -----  
# Random Forest  
# -----  
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)  
rf_model.fit(X_train, y_train)  
y_pred_rf = rf_model.predict(X_test)  
  
print("\n📊 Random Forest Results:")  
print("Accuracy:", accuracy_score(y_test, y_pred_rf))  
print(classification_report(y_test, y_pred_rf))  
sns.heatmap(confusion_matrix(y_test, y_pred_rf), annot=True, fmt='d', cmap='Greens')  
plt.title("Random Forest - Confusion Matrix")  
plt.show()
```

 <ipython-input-17-3fb8eed057d3>:17: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series, and this may result in modifications to the original data. The behavior will change in pandas 3.0. This inplace method will never work because the original data is not a copy.  
For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value})'.


```
df['TotalCharges'].fillna(df['TotalCharges'].median(), inplace=True)
```

 Logistic Regression Results:

Accuracy: 0.8211497515968772

	precision	recall	f1-score	support
0	0.86	0.90	0.88	1036
1	0.69	0.60	0.64	373
accuracy			0.82	1409
macro avg	0.77	0.75	0.76	1409
weighted avg	0.82	0.82	0.82	1409


Logistic Regression - Confusion Matrix



 Random Forest Results:

Accuracy: 0.801277501774308

	precision	recall	f1-score	support
0	0.83	0.91	0.87	1036
1	0.67	0.49	0.57	373
accuracy			0.80	1409
macro avg	0.75	0.70	0.72	1409
weighted avg	0.79	0.80	0.79	1409



## ✓ Model Evaluation

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import (
    accuracy_score, f1_score, roc_auc_score,
    mean_squared_error, confusion_matrix, roc_curve
)

# Load dataset
df = pd.read_csv("Telco-Customer-Churn.csv")

# Data preprocessing
df.drop('customerID', axis=1, inplace=True)
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], errors='coerce')
```

```
df['TotalCharges'].fillna(df['TotalCharges'].median(), inplace=True)

le = LabelEncoder()
df['Churn'] = le.fit_transform(df['Churn'])
for col in df.select_dtypes(include='object').columns:
    if df[col].nunique() == 2:
        df[col] = le.fit_transform(df[col])
    else:
        df = pd.get_dummies(df, columns=[col])

scaler = StandardScaler()
df[['tenure', 'MonthlyCharges', 'TotalCharges']] = scaler.fit_transform(df[['tenure', 'Monthl

# Split data
X = df.drop("Churn", axis=1)
y = df["Churn"]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train models
log_model = LogisticRegression(max_iter=1000)
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
log_model.fit(X_train, y_train)
rf_model.fit(X_train, y_train)

# Predictions
log_preds = log_model.predict(X_test)
rf_preds = rf_model.predict(X_test)

# Evaluation metrics
print("Logistic Regression")
print("Accuracy:", accuracy_score(y_test, log_preds))
print("F1 Score:", f1_score(y_test, log_preds))
print("ROC AUC:", roc_auc_score(y_test, log_model.predict_proba(X_test)[: , 1]))
print("RMSE:", np.sqrt(mean_squared_error(y_test, log_preds)))
print(confusion_matrix(y_test, log_preds))

print("\nRandom Forest")
print("Accuracy:", accuracy_score(y_test, rf_preds))
print("F1 Score:", f1_score(y_test, rf_preds))
print("ROC AUC:", roc_auc_score(y_test, rf_model.predict_proba(X_test)[: , 1]))
print("RMSE:", np.sqrt(mean_squared_error(y_test, rf_preds)))
print(confusion_matrix(y_test, rf_preds))

# Confusion matrix plots
sns.heatmap(confusion_matrix(y_test, log_preds), annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix - Logistic Regression")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()


sns.heatmap(confusion_matrix(y_test, rf_preds), annot=True, fmt='d', cmap='Greens')
```

```
plt.title("Confusion Matrix - Random Forest")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

# ROC curves
log_fpr, log_tpr, _ = roc_curve(y_test, log_model.predict_proba(X_test)[:, 1])
rf_fpr, rf_tpr, _ = roc_curve(y_test, rf_model.predict_proba(X_test)[:, 1])

plt.figure(figsize=(8, 6))
plt.plot(log_fpr, log_tpr, label='Logistic Regression')
plt.plot(rf_fpr, rf_tpr, label='Random Forest')
plt.plot([0, 1], [0, 1], 'k--')
plt.title("ROC Curve")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend()
plt.grid(True)
plt.show()
```



 <ipython-input-18-84bfb83270b3>:20: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series resulting in 2D indexing. The behavior will change in pandas 3.0. This inplace method will never work because the operation is not performed on the original data. For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value})'.

```
df['TotalCharges'].fillna(df['TotalCharges'].median(), inplace=True)
```

Logistic Regression

Accuracy: 0.8211497515968772

F1 Score: 0.64

ROC AUC: 0.8621127351020114

RMSE: 0.42290690276126114

```
[[933 103]
```

```
 [149 224]]
```

Random Forest

Accuracy: 0.801277501774308

F1 Score: 0.56656346749226


ROC AUC: 0.8385313693624687

RMSE: 0.4457830169776457

```
[[946  90]
```

```
 [190 183]]
```

Confusion Matrix - Logistic Regression



Confusion Matrix - Random Forest



